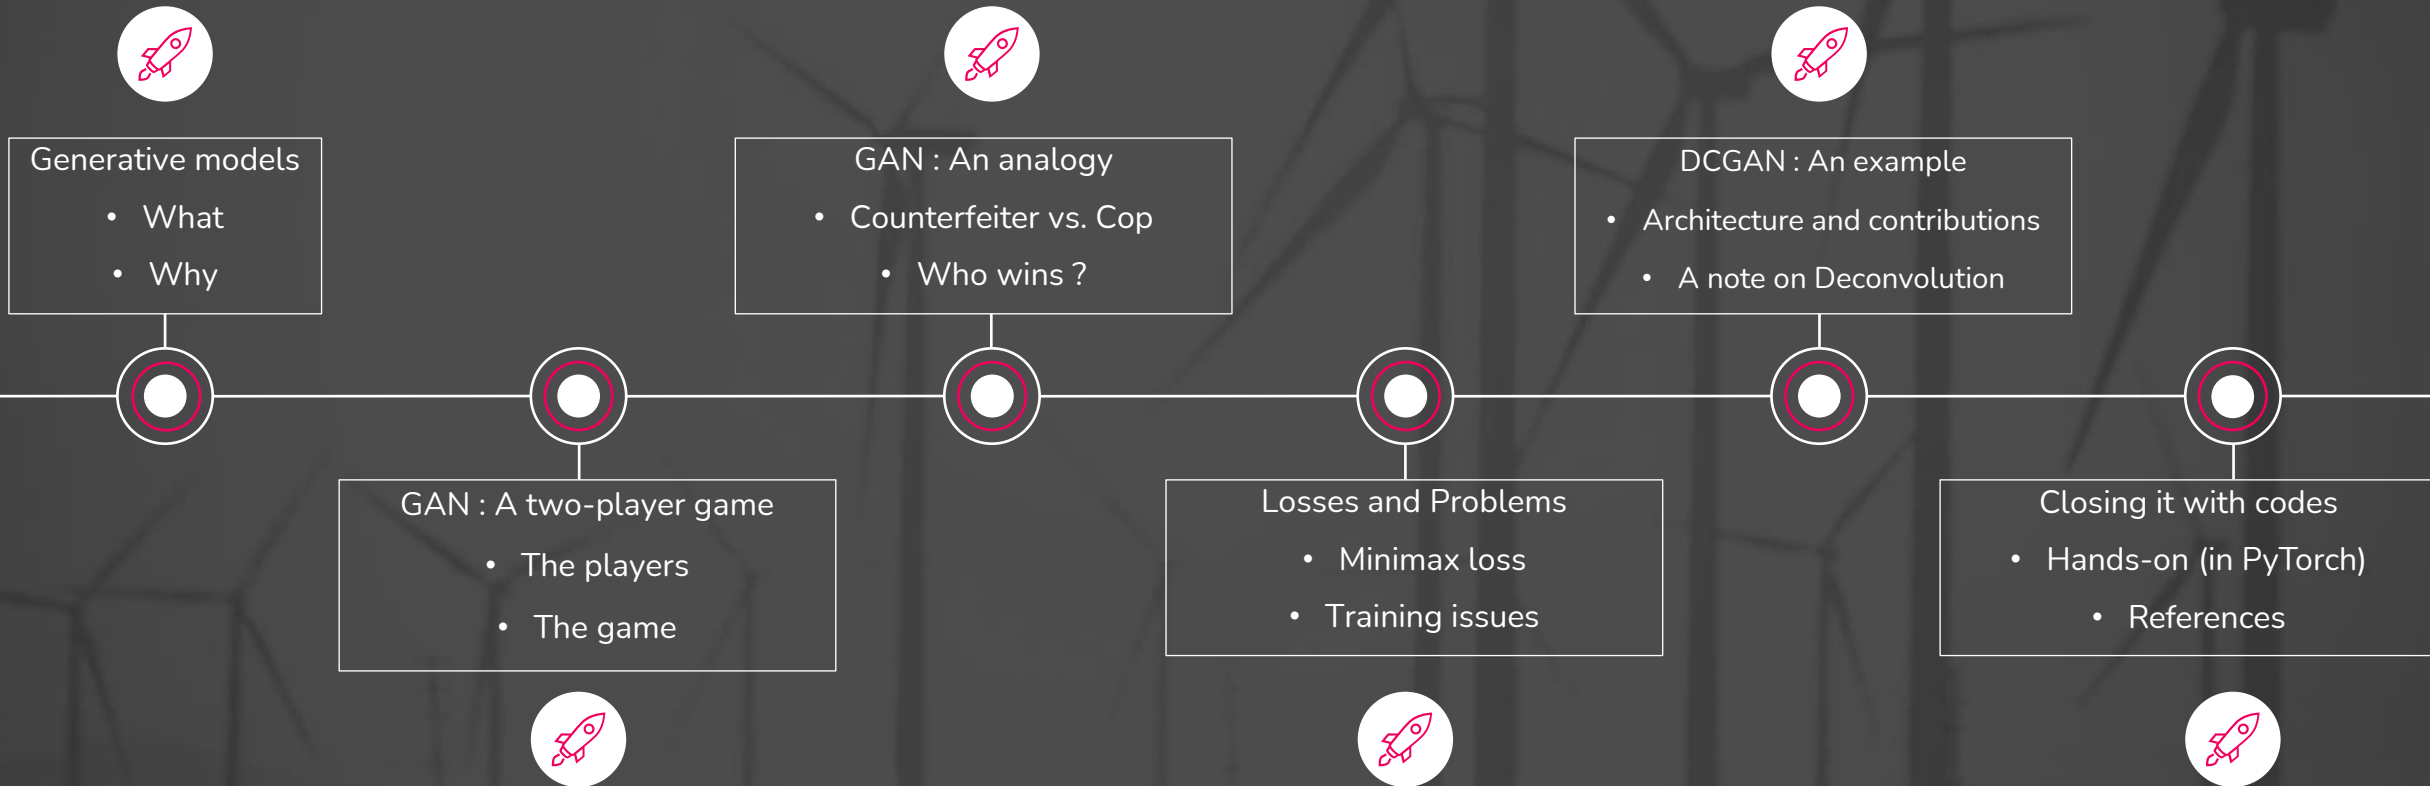


Generative Adversarial Networks

A CS590 Class



What are generative models? What's the need?

Create new data instances that look like the training data !

Attribute2Font



Anime face generation

More applications!




Image de-raining



Image super-resolution



Text2Image



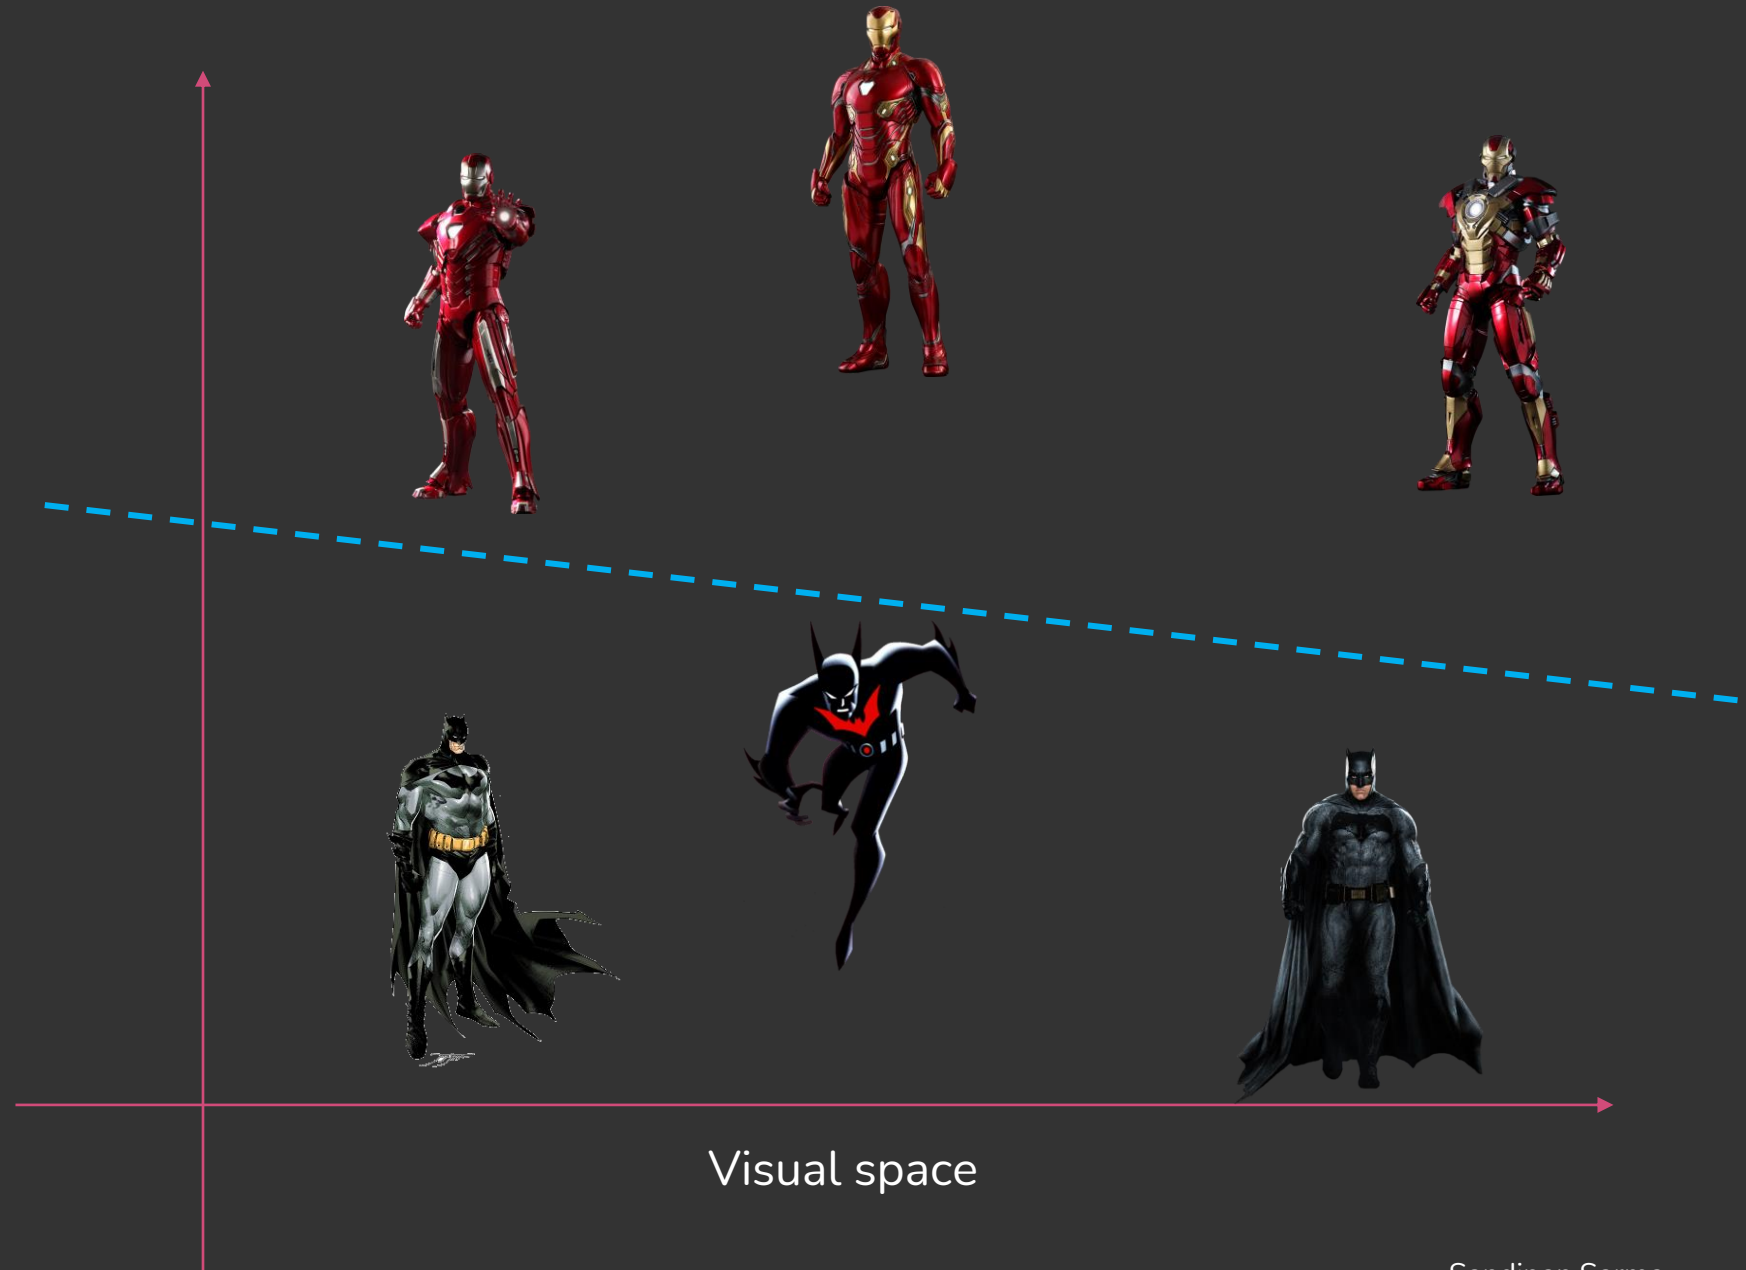
Face completion

“GAN is a two-player minimax game”

Let's first familiarize ourselves with the players and their purpose.....

Purpose of Player A

- Look at different **Batsuits** and **Ironman armors** (y) and corresponding visual images (x)
 - Mathematically, learn $P(y|x)$
 - Learn **difference** between Ironman and Batman
-
- Hence, player name:
Discriminative model




Purpose of Player B

- Look at different **Batsuits** and **Ironman armors** (y) and corresponding visual images (x)
- Mathematically, learn $P(x,y)$
- Learn to imitate data distribution of Ironman and Batman
- **Generate** new suits from each distribution
- Hence, player name:

**Generative
model**

IIT Guwahati
October 18, 2022



Sandipan Sarma
Research Scholar

Counterfeiting currency bills



**Tries to generate real-like
counterfeit bills**



**Tries to catch fake bills, but
incurs penalty if he fails**

An amateur counterfeiter gets caught



**Initial attempt at
counterfeiting bills**



Fake bills easily identified

Counterfeiter learns from previous mistakes and successfully fools the cop



Improvement to generate
even better bills



You're fired.

Cannot spot the difference between real
and fake note this time, so incurs a loss
(penalty)

A better cop can still identify the fake bills



Improvement to generate even better bills

Understands discrimination between real and fake notes in greater detail

Further improvement by counterfieter



Achieving expertise at counterfeiting

(GENERATOR)



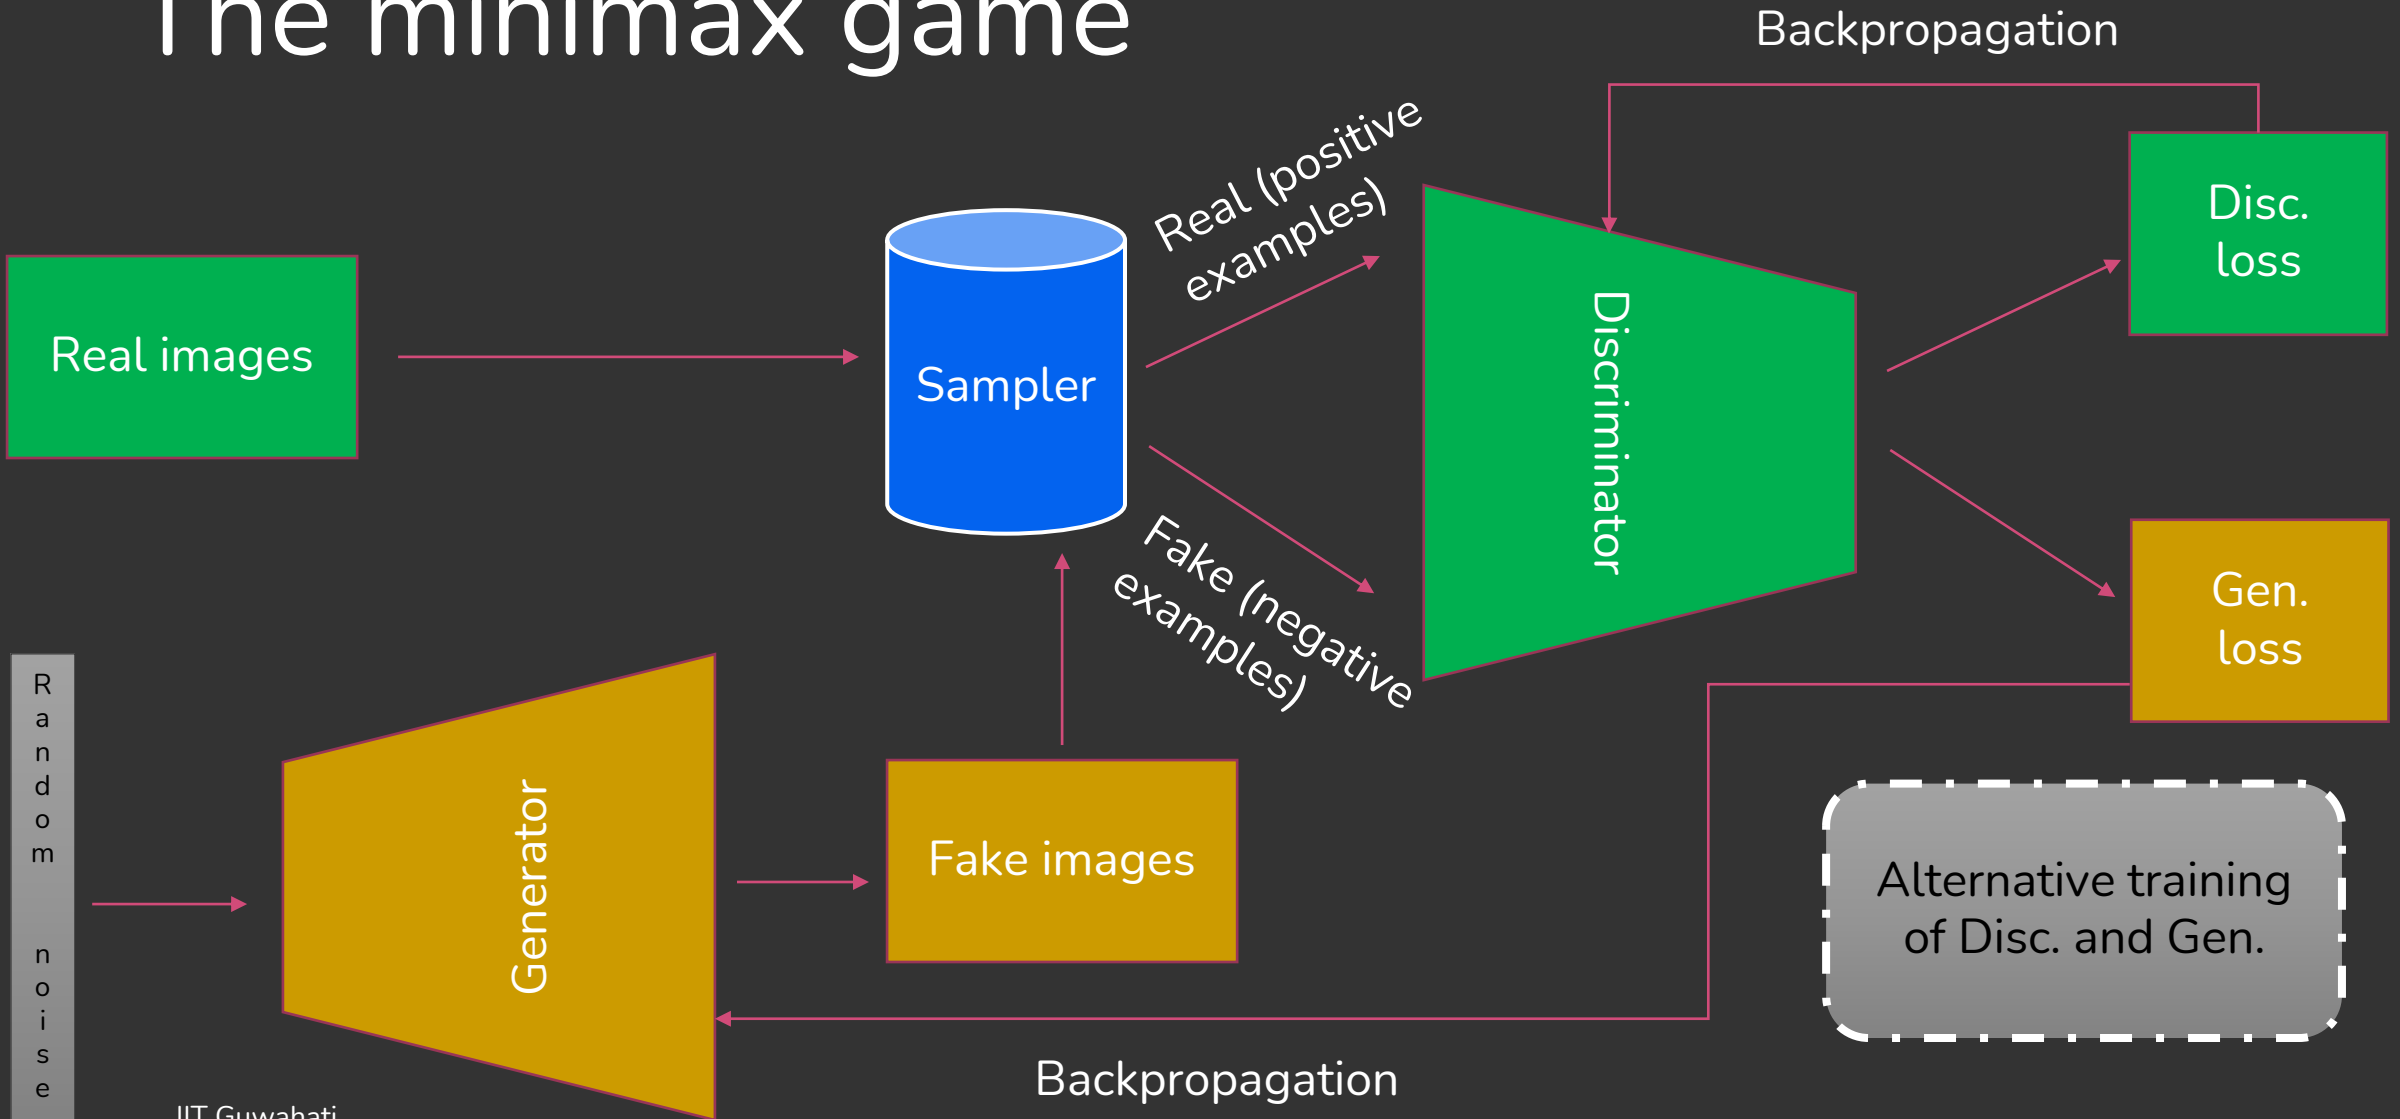
Looks almost real !

(DISCRIMINATOR)

Winner?

Both improve
themselves through
each other's feedback

The minimax game



Why differentiability?

- Most machine learning algorithms formulated as **convex optimization problems**
- Gradient computations
- Differentiable functions don't have sudden jerks – can be **interpolated/extrapolated** to get predicted value reasonably close to actual value

Loss functions

Minimax loss (original paper on GAN [7]):

$$E_x[\log(D(x))] + E_z[\log(1 - D(G(z)))]$$


Other loss functions like WGAN-loss, CGAN loss, InfoGAN loss, etc. are subject to the GAN architecture in general

Some problems faced during training

- Vanishing gradients
- Failure in convergence
- Unstable training
- Mode collapse
- Lack of proper evaluation metric to check GAN progress – evolution of loss does not necessarily indicate the reach of equilibrium

Deep Convolutional GAN (DCGAN) [8]

Model architecture




Some important contributions of DCGAN


- Replace pooling layers of discriminator (D) with strided conv. layers and those of generator (G) by FRACTIONALLY-STRIDED CONV. layers
- Eliminate fully-connected layers at the end of model
- Use BatchNorm for both G and D
- Use ReLU in G for all layers except last layer, which uses *tanh*
- Use LeakyReLU in D for all layers

Some useful hacks for training a DCGAN (as suggested by its authors [8]) can be found at : [Ganhacks](#)

Quick note: Fractionally-strided convolution / Transpose convolution



Input = (2,2)
Stride = (1,1)
Padding = None
Output = (4,4)



Input = (3,3)
Stride = (2,2)
Padding = Valid
Output = (5,5)

- See detailed discussion at <https://distill.pub/2016/dec-conv-checkerboard/>
- Find more CONV visualizations at [Conv_arithmetic](#)

Towards implementing DCGAN

- Deep learning framework used : PyTorch
- Dataset : [Celeb-A](#)
- Link to code notebook : [DCGAN-tutorial](#)
- Outline:
 - ✓ Basic PyTorch pre-requisites
 - ✓ Setup parameters
 - ✓ Data
 - ✓ Weight initializers
 - ✓ Generator and Discriminator architectures
 - ✓ Loss function and optimizer
 - ✓ Training Generator and Discriminator
 - ✓ Results
 - ✓ What can you explore?

References

- [1] Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2019.
- [2] Zhang, Han, et al. "Stackgan: Text to photo-realistic image synthesis with stacked generative adversarial networks." *Proceedings of the IEEE international conference on computer vision*. 2017.
- [3] Wang, Yizhi, Yue Gao, and Zhouhui Lian. "Attribute2Font: creating fonts you want from attributes." *ACM Transactions on Graphics (TOG)* 39.4 (2020): 69-1.
- [4] Zhang, He, Vishwanath Sindagi, and Vishal M. Patel. "Image de-raining using a conditional generative adversarial network." *IEEE transactions on circuits and systems for video technology* 30.11 (2019): 3943-3956.
- [5] Ledig, Christian, et al. "Photo-realistic single image super-resolution using a generative adversarial network." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.
- [6] Li, Yijun, et al. "Generative face completion." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.
- [7] Goodfellow, Ian, et al. "Generative adversarial nets." *Advances in neural information processing systems* 27 (2014).
- [8] Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks." *arXiv preprint arXiv:1511.06434* (2015).